



OpenMM workshop





Welcome and Introductions

Vijay Pande
OpenMM Workshop



Introductions

- **Stanford/Simbios team development**

- Peter Eastman (OpenMM API, GPU code)
- Yutong Zhou (GPU code, Folding@home)
- Lee-Ping Wang, Diwakar Shukla, Gert Kiss (Simbios Distinguished Fellows)
- Vijay Pande (co-PI, protein folding DBP lead PI)

- **Collaborators**

- Kyle Beauchamp (PyMD)
- John Chodera (Yank, PyOpenMM, tests)
- Erik Lindahl (Stockholm)

- **Simbios**

- Joy Ku (Director of Dissemination)
- Russ Altman & Scott Delp (Simbios co-PI's)



Introduction of OpenMM

Vijay Pande
OpenMM Workshop



OpenMM



= rapid development +
rapid execution

OpenMM is an app, API, and library for rapid
molecular dynamics.

Easy to modify and incorporate into any code.

Two competing programming challenges


- Rapid execution, but slow development
 - traditional model
 - **Examples:** C, FORTRAN, Assembly
- Rapid development, but slow execution
 - new languages often used by scientists
 - **Examples:** Python, Perl
- We can have it both ways: Domain specific languages (DSLs)
 - New paradigm: specific language for application area
 - **Examples:** OpenGL, SQL, Matlab
- A DSL approach for molecular mechanics?
 - OpenMM hopes to provide **both** rapid execution and rapid development

Goals

- **A complete library for molecular mechanics**
 - complete = what would need to do to do the most common calculations
 - complete != does everything conceivable
- **Fast and general**
 - **“rapid development *and* rapid execution”**
 - Don't exposure hardware specifics
 - but optimize for speed underneath
 - broad support for GPUs, multicore, etc
- **Two level API**
 - OpenMM for high level: intended for application developers
 - Low level API: for developers (mainly in-house & accelerator devs)
 - OpenMM can be a nexus for application and low level programmers to meet

GPUs: very useful for rapid execution

- **Graphics Processing Units (GPUs) are very powerful**
 - Folding@home calculation circa 2003 = 10,000 PC's @ 1 GFLOP/PC = 10,000 GFLOPS
 - Fast GPU today = 1,000 GFLOP
 - Fast GPU cluster today = ~50,000 GFLOPS
- **GPU's are getting faster, faster than CPU's**
 - Moore's law is dead for traditional CPU's
 - we now see more cores per chip, but each core isn't any faster
 - GPU's figured out this trick a long time ago
 - typical GPU's now have 100's of cores
 - GPU's use their cores more efficiently
- **BUT, GPU's are horrible to program**
 - can't just recompile
 - must rethink algorithms
 - must understand the nature of the hardware
 - work closely with vendors (we collaborate closely with AMD/ATI, NVIDIA, and Intel)



ATI X1900XT (500 GFlops peak, ~\$100 + of a cost computer)



Sony PS3 (Cell processor: 220 GFlops peak, ~\$400 total)

Unique aspects of comp biology on GPUs

- **Design algorithms that are GPU friendly**

- FLOPS are free, memory is expensive
- low lying fruit: algorithms which map well to GPUs

- **Code everything on the GPU**

- If the original bottleneck is 90% of the calc, that's still only a 10x speed up at best
- to get 100x to 1000x, one needs to have the whole calculation on the GPU (in our experience)

- **Centralized libraries, open source (eg OpenMM)**

- avoid reinventing the wheel
- build on others' work



- **Next steps**

- not just speeding existing algorithms, but new methods
- code methods which we wouldn't even dare to try now

The opportunity for a common library

- **The molecular mechanics community has become fragmented**
 - tens of different MD codes
 - hardware acceleration (via multi-core, SSE, MPI, GPU's, and math coprocessors) is a critical element
 - much like the graphics community in the 1980's
- **Hardware acceleration is a great unifying factor**
 - unifying API the way OpenGL unified graphics
 - incorporates hardware acceleration in its base design
 - this API would be used as the backend to existing codes, allowing for all to benefit from hardware acceleration

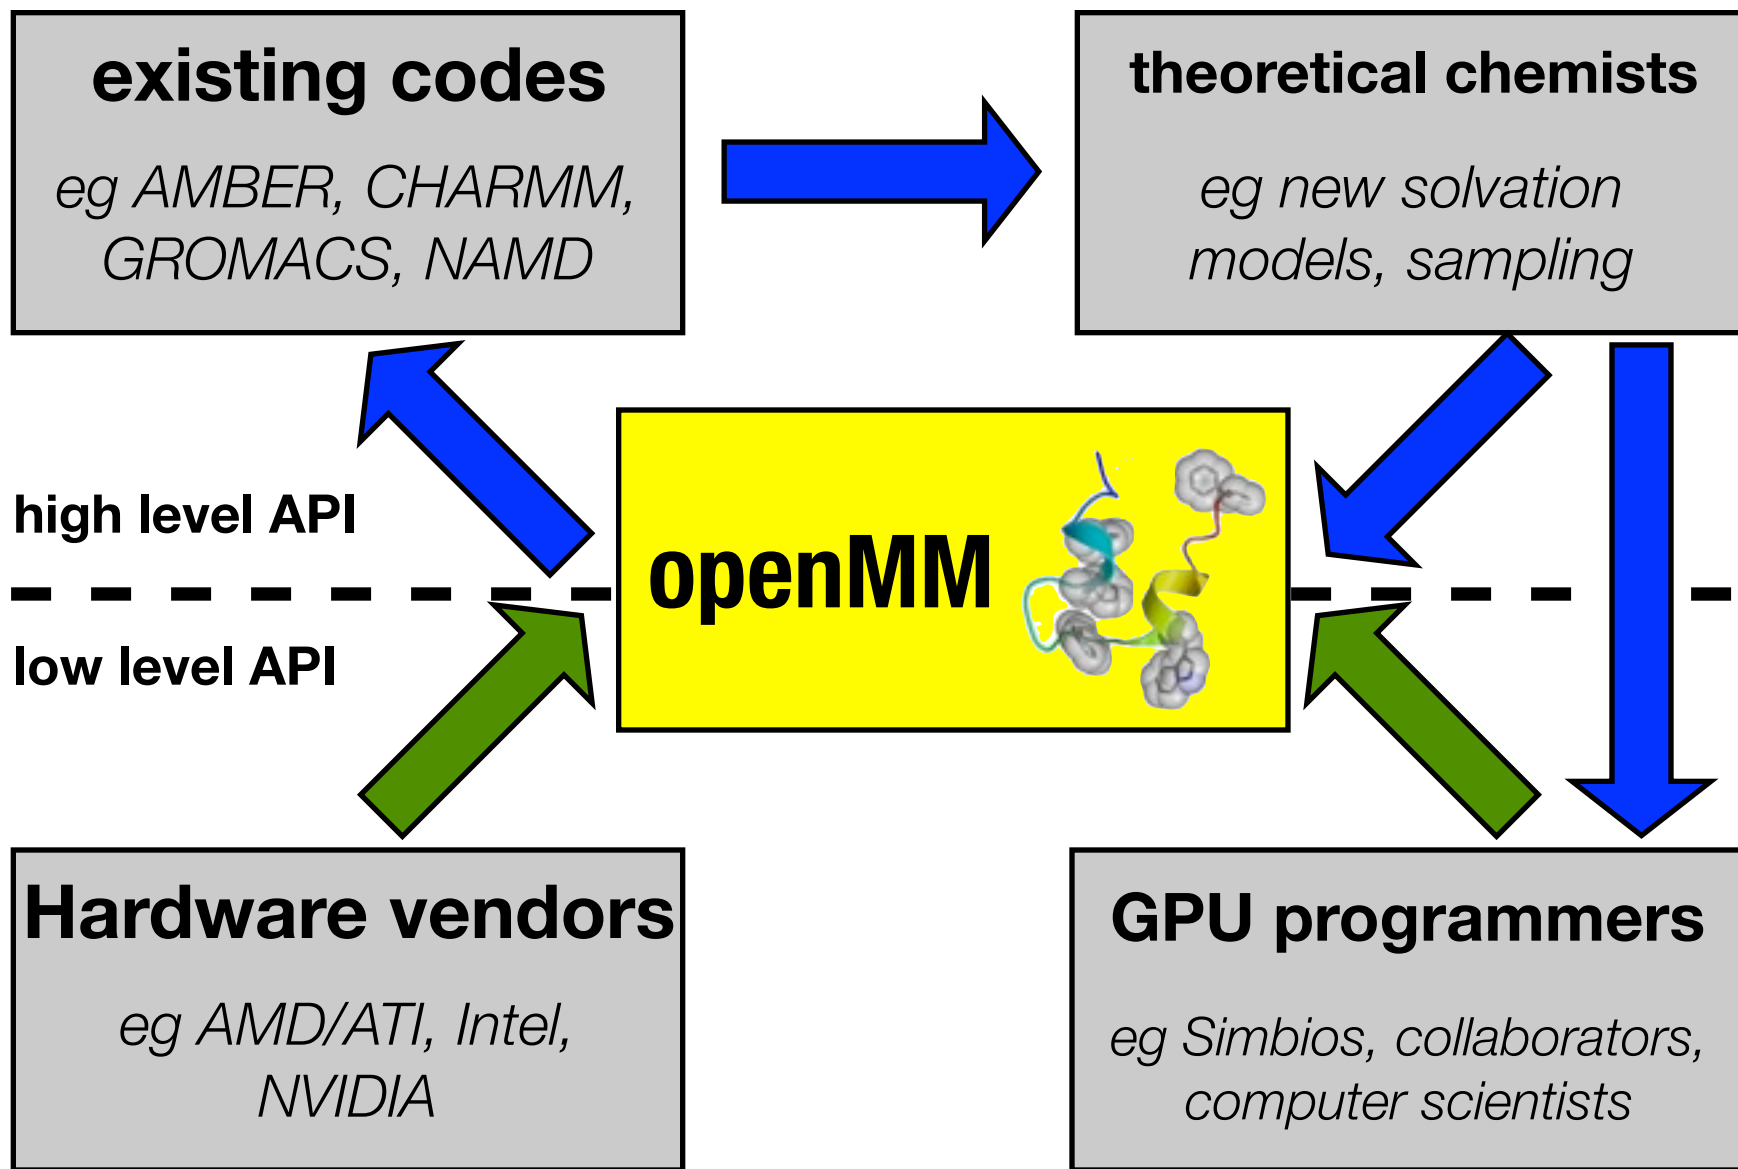


ATI X1900XT (500 GFlops peak, ~\$100 + of a cost computer)

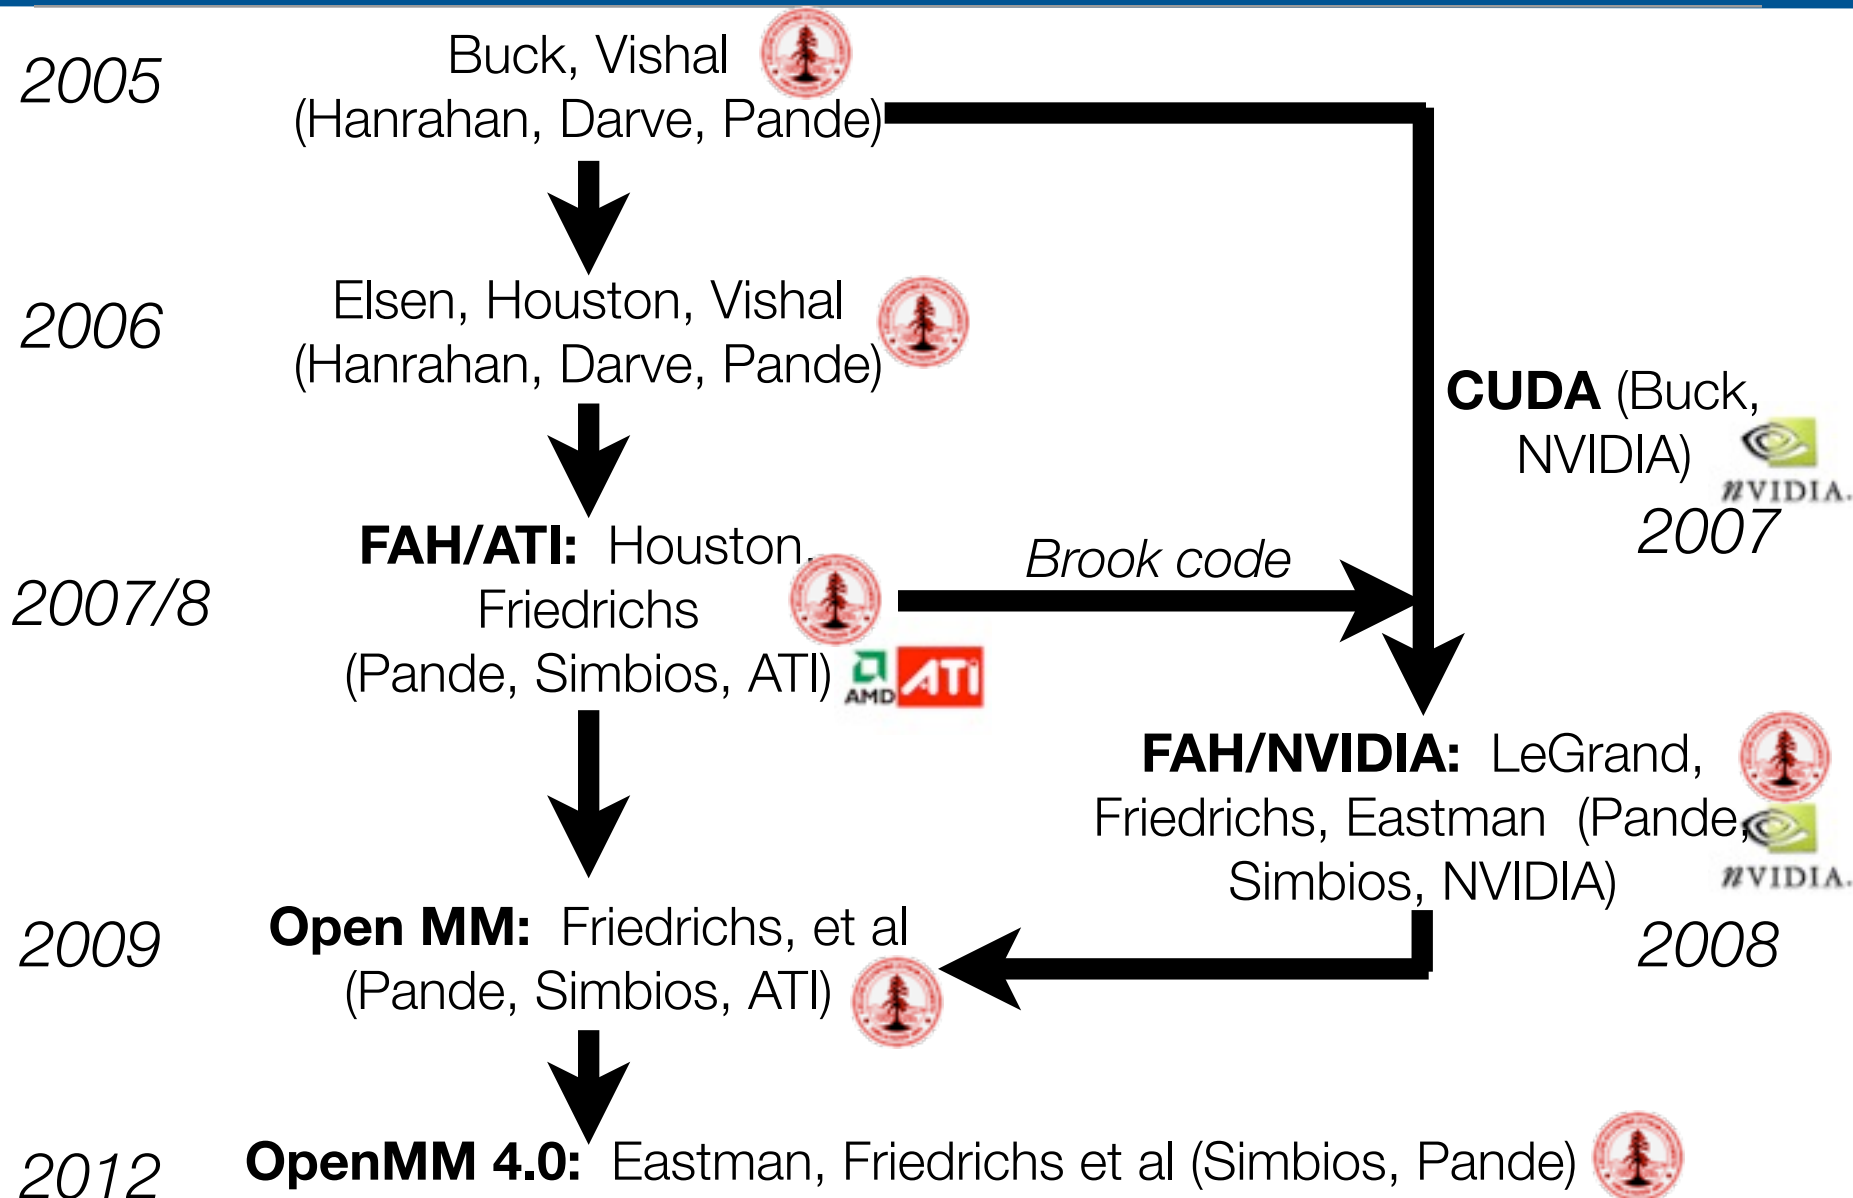


Sony PS3 (Cell processor: 220 GFlops peak, ~\$400 total)

Connections to OpenMM



History of OpenMM



OpenMM 5.1 speed improvements: JAC benchmark

	CUDA (GTX 680)	OpenCL (GTX 680)	OpenCL (HD 7970)
Implicit hbonds	148	134	120
Implicit hangles	240	209	104
RF hbonds	90.3	78.1	83.5
RF hangles	141.1	113.0	90.2
PME hbonds	61.0	41.5	49.3
PME hangles	99.9	66.9	63.0

Joint AMBER-CHARMM DHFR Benchmark in ns/day

OpenMM is under active development

- 4.1 (~April 2012): Modeling tools, bugfixes
- 4.0 (January 2012): Python based app, multi-GPU support, Ring Polymer MD (RPMD)
- 3.1.1 (Aug. 11, 2011): bug fixes
- 3.1 (Aug. 1, 2011): performance improvements, especially to the AMOEBA plugin; support for parallelizing computations across multiple GPUs; support of direct polarization model by AMOEBA plugin; GB/VI force fully supported
- 3.0 (Mar. 30, 2011): supports AMOEBA force field; provides an energy minimizer; CMAP torsions; improved performance, especially for running on CPUs; Python API wrappers
- 2.0 (June 24, 2010): supports pressure coupling; provides custom Hbond forces; major performance improvements; works on ATI GPUs
- 1.1.1 (Mar. 4, 2010): bug fixes
- 1.1 (Feb. 12, 2010): provides custom torsion forces; bug fixes and other improvements
- 1.0 (Jan. 20, 2010): provides new custom forces, including bonds, angles, external, and GB; improved OpenCL support
- 1.0 Beta (Oct. 30, 2009): supports Particle Mesh Ewald (PME); custom nonbonded interactions, including algebraic and tabulated forces; preliminary OpenCL support
- Preview Release 4 (Aug. 20, 2009): supports Ewald summation; GPU accelerated energy calculations; C and Fortran API wrappers; a faster constraint algorithm; and many minor enhancements
- Preview Release 3 (May 19, 2009): provides support for explicit solvent on NVIDIA GPUs; includes periodic boundary conditions, cutoffs on non-bonded interactions, and new constraint algorithms.
- Preview Release 2 (Jan. 26, 2009): provides support for accelerating molecular modeling simulations on ATI and NVIDIA graphics processor units (GPUs)

OpenMM roadmap: 2013

- **OpenMM 5.1**

- Big speed increases over 5.0 (2x faster)
- Integration with CHARMM
- Accelerated MD

- **OpenMM 5.2**

- ABSINTH implicit solvent
- AMOEBA OpenCL implementation
- Colored noise integrator
- LTMD
- CHARMM* force fields

- **Longer term**

- CUDA JIT (when CUDA supports it)
- More optimizations
- Custom compound bonded force
- Finalize Rosetta force field
- Triclinic boxes
- A more accurate SASA calculation for use with GB models
- QM/MM

<http://wiki.simtk.org/openmm/RoadmapTimeline>

Looking further to the future: FAQ

- **Scheme to manage API for years to come**

- design decisions were made to address trade-off between generality/performance: goals are performance and simplicity, features come 2nd
- speed/simplicity/elegance were the motivations for these decisions

- **How will the API be updated in the future?**

- As OpenMM has progressed, we have brought in other key developers to advise us, including Erik Lindahl (Gromacs), Jesus Izaguirre (Protomol), & Ross Walker (AMBER) as well as the Simbios SAB to advise us for useful directions for the API
- We plan to formalize this to have an advisory board for OpenMM specifically


- **How difficult is it for hardware vendors to produce mature and performant OpenMM “drivers”?**

- Strong success working with NVIDIA, ATI, and Sony and are making connections/collaborations with other key players (OpenCL, Larrabee)
- Due to its design and the existence of key OpenMM apps (eg Folding@home, Gromacs), OpenMM will be able to stay current

Examples of current OpenMM enabled applications



Folding@home
<http://folding.stanford.edu>




Gromacs
<http://www.gromacs.org>

CHARMM

CHARMM

<http://www.charmm.org>



NAST

<http://simtk.org/home/nast>



Protomol

<http://protomol.sourceforge.net>



Yank

<http://simtk.org/home/yank>

Licensing and distribution

- **API & reference BSD license, GPU kernels are LGPL**
 - free & open
 - we want LGPL to have a community owned set of GPU kernels
 - we're looking for collaborations for new features
- **But, please cite us**
 - Any work that uses OpenMM should cite the following paper:
M. S. Friedrichs, P. Eastman, V. Vaidyanathan, M. Houston, S. LeGrand, A. L. Beberg, D. L. Ensign, C. M. Bruns, V. S. Pande.
“Accelerating Molecular Dynamic Simulation on Graphics Processing Units.” J. Comp. Chem., (2009)
 - [http://www3.interscience.wiley.com/journal/121677402/
abstract](http://www3.interscience.wiley.com/journal/121677402/abstract)

Summary

- **What is it**

- API, library, and application for core molecular dynamics / molecular mechanics applications
- emphasis on rapid development *and* rapid execution speed (eg hardware acceleration)
- dual APIs (one for applications and one for low level hardware)
- open source (LGPL) software

- **What is it not**

- a general solution for all possible molecular mechanics tasks
- a compiler which can turn a generic MD code into accelerated code
